

# Computer Vision - Image Features

Włodzimierz Kasprzak and Artur Wilkowski

26 November 2021

- Feature: “an interesting part of an image”
- Main Goal: Make compute vision problems tractable by reducing full-size image to a limited set of features
- Two subproblems: Feature detection and feature description

We will discuss:

- Corner detection
- Scale space features
- Geometric primitives

## Feature detection: Corner features

## Moravec operator (1977)

The detector is based on the auto-correlation function of the signal:

- It measures the grey value differences between a window and windows shifted in several directions;
- Four discrete shifts in directions parallel to the rows and columns of the image are used.

PROCEDURE outline:

- 1 FOR every pixel  $(x, y)$  compute average differences of the image function along four directions:

$$V_0 = \frac{1}{20} \sum_{i=-2}^2 \sum_{j=-2}^1 [f(x+i, y+j) - f(x+i, y+j+1)]^2$$

$$V_1 = \frac{1}{20} \sum_{i=-2}^1 \sum_{j=-2}^2 [f(x+i, y+j) - f(x+i+1, y+j)]^2$$

$$V_2 = \frac{1}{16} \sum_{i=-2}^1 \sum_{j=-2}^1 [f(x+i, y+j) - f(x+i+1, y+j+1)]^2$$

$$V_3 = \frac{1}{16} \sum_{i=-2}^1 \sum_{j=-2}^1 [f(x+i, y+j) - f(x+i+1, y+j-1)]^2$$

- 2 Then for every pixel  $(x,y)$  compute

$$M = \min\{V_0, V_1, V_2, V_3\}$$

- 3 if  $M > \theta$  set pixel  $(x,y)$  as a keypoint

# Harris-Stephens corner detector

Harris-Stephens detector improves on Moravec detector by considering every not only discrete 45deg. corner rotations

PROCEDURE outline:

FOR every pixel  $(x, y)$ :

- 1 Create and auto-correlation matrix  $A(x, y)$  that averages discrete derivatives of the image in a window  $W(x, y)$  around a point  $P(x, y)$ :

$$A(x, y) = \sum_{(x_k, y_k) \in W} w(x - x_k, y - y_k) \begin{pmatrix} I_x^2(x_k, y_k) & I_x(x_k, y_k)I_y(x_k, y_k) \\ I_x(x_k, y_k)I_y(x_k, y_k) & I_y^2(x_k, y_k) \end{pmatrix}$$

with  $w$  being the windowing function (gaussian for best results)

# Harris-Stephens corner detector

PROCEDURE outline:

Perform Gaussian smoothing of the image then FOR every pixel ( $x$ ,  $y$ ):

- 1 Create an auto-correlation matrix  $A(x, y)$  that averages discrete derivatives of the image in a window  $W(x, y)$  around a point  $P(x, y)$ :

$$A(x, y) = \sum_{(x_k, y_k) \in W} w(x - x_k, y - y_k) \begin{pmatrix} I_x^2(x_k, y_k) & I_x(x_k, y_k)I_y(x_k, y_k) \\ I_x(x_k, y_k)I_y(x_k, y_k) & I_y^2(x_k, y_k) \end{pmatrix}$$

- 2 Eigenvalues  $\lambda_1, \lambda_2$  of  $A$  are to be considered (not necessarily explicitly evaluated)
- 3 Corner response (using constant  $k$  is)

$$R = \lambda_1 \lambda_2 - k(\lambda_1 + \lambda_2)^2 = \det(A) - k(\text{trace}(A))^2$$

and  $k$  is a small constant

- 4 Corners are selected after application of non-max suppression procedure



Shi-Thomasi Good Features to Track algorithm is essentially a Harris corner detector with a modified corner response function yielding better results

$$R = \min(\lambda_1, \lambda_2)$$

# Multi-scale features

# Scale problem of feature detection

- Problem: features in the image can have arbitrary size - how to find the best scale for a feature
- Solution: perform search both in space and scale (scale-space search)

## Laptev Lindberg classic scale-space representation

- Scalespace is obtained by convolution of the original image with a Gaussian kernel for different values of  $\sigma$

$$L(\cdot, s) = g(\cdot, s) * f$$

- The Gaussian kernel is given by

$$g(x, y, s) = \frac{1}{2\pi s} e^{-(x^2+y^2)/(2s)}$$

- In order to detect features we use derivatives of the scale space  $L$

$$L_{x^\alpha, y^\alpha}(\cdot, t) = \partial_{x^\alpha, y^\alpha} L(\cdot, t) = g_{x^\alpha, y^\alpha}(\cdot, t) * f$$

- In order for the derivatives to be comparable between scales, they are normalized

$$\partial_\xi = \sqrt{t} \partial_x, \text{ or } \partial_{x, \gamma\text{-norm}} = t^{\gamma/2} \partial_x$$

# Scale space representation

- Detection of blobs is achieved by looking up scale space maxima for scaled Laplacian-of-Gaussians

$$\Delta^2 L = t(L_{xx} + L_{yy})$$

- Detection of elongated ridges is achieved by:
  - Computation of the Hessian matrix

$$H = s^\gamma \begin{pmatrix} L_{xx} & L_{xy} \\ L_{xy} & L_{yy} \end{pmatrix}$$

- Computation of the eigenvalues of  $H$

$$L_{pp,\gamma-norm} = \frac{s^\gamma}{2} \left( L_{xx} + L_{yy} - \sqrt{(L_{xx} - L_{yy})^2 + 4L_{xy}^2} \right)$$





$$L_{qq,\gamma-norm} = \frac{s^\gamma}{2} \left( L_{xx} + L_{yy} + \sqrt{(L_{xx} - L_{yy})^2 + 4L_{xy}^2} \right)$$

- Comparison of the eigenvalues





$$\mathcal{A}_{\gamma-norm} L = (L_{pp,\gamma-norm} - L_{qq,\gamma-norm})^2$$

- Selection of scale-space extrema according to  $\mathcal{A}$


# Scale space representation



# Scale space representation



# Scale space representation





## Scale-Invariant Feature Transform

- Detection of optimum keypoints in scale-space
- Computed features have different sizes
- Detected features have are described by a triplet  $(\mathbf{x}_R, s, \alpha)$

## PROCEDURE outline

- 1 Compute scale-space by convolving image with Gaussian kernels and subsampling between octaves

- Image convolution

$$L(x, y, \sigma) = G(x, y, \sigma) * I(x, y)$$

with a kernel


$$G(x, y, \sigma) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2}$$

- Computation of the Difference of Gaussians

$$\begin{aligned} D(x, y, \sigma) &= (G(x, y, k\sigma) - G(x, y, \sigma)) * I(x, y) \\ &= L(x, y, k\sigma) - L(x, y, \sigma) \end{aligned}$$

- The DoG is close approximation to the LoG


# Scale-Invariant Feature Transform



Source: David. G. Lowe, Distinctive image features from scale-invariant keypoints, 2004

# Scale-Invariant Feature Transform

- 2 Perform non-extrema suppression using neighbours from the current image and from one scale up and down
  - Image convolution



Source: David. G. Lowe, Distinctive image features from scale-invariant keypoints, 2004

- 3 Accurate keypoint localization. Sub-pixel approximation of the local extremum is obtained by applying the Taylor series expansion (up to quadratic terms)

$$D(\mathbf{x}) = D + \frac{dD^T}{d\mathbf{x}}\mathbf{x} + \frac{1}{2}\mathbf{x}^T \frac{d^2D}{d\mathbf{x}^2}\mathbf{x}$$

And the optimum is given by

$$\hat{\mathbf{x}} = \frac{d^2D^{-1}}{d\mathbf{x}^2} \frac{dD}{d\mathbf{x}}$$

4. Weak extrema are eliminated by thresholding  $|D(\hat{\mathbf{x}})| < 0.03$  ( $D$  is normalized within  $[0, 1]$ )
5. Edges are eliminated from the set of detected extrema by analyzing eigenvalues of Hessian Matrix computed as

$$H = \begin{pmatrix} D_{xx} & D_{xy} \\ D_{xy} & D_{yy} \end{pmatrix}$$

Extrema with a small one eigenvalue and a large second eigenvalue are rejected as edges. As in Harris-Sepens operator, the eigenvalues do not need to be explicitly computed.

## 6. Gradient directions and magnitudes are computed

$$m(x, y) = \sqrt{(L(x+1, y) - L(x-1, y))^2 + (L(x, y+1) - L(x, y-1))^2}$$
$$\theta(x, y) = \text{atan2}(L(x, y+1) - L(x, y-1), L(x+1, y) - L(x-1, y))$$

## 7. Establishing keypoint orientation by the gradient directions and magnitudes in keypoint neighbourhood.

- 36 orientation bins is used
- each gradient votes with its magnitude smoothed by the Gaussian kernel
- the highest histogram peak together with all peaks above 80% of the highest one are selected
- in effect: a keypoint can be multiplied with different orientations

# Scale-Invariant Feature Transform



Source: [https://docs.opencv.org/3.0-beta/doc/py\\_tutorials/py\\_feature2d/py\\_sift\\_intro/py\\_sift\\_intro.html](https://docs.opencv.org/3.0-beta/doc/py_tutorials/py_feature2d/py_sift_intro/py_sift_intro.html)




## Feature detection: Geometric primitives

# Approximating contour with a polyline

- The Douglas-Peucker algorithm tries to approximate a contour of arbitrary length using a small number of linear segments, by splitting initial segment
- Splitting stops when all points corresponding to the line segment are closer to the segment than some threshold value  $\theta$
- Distance from the line segment  $a \cdot x + b \cdot y + c = 0$  to some point  $(x_i, y_i)$  is given by

$$d_i = \frac{|ax_i + by_i + c|}{a^2 + b^2}$$



# Approximating contour with a polyline

## Douglas–Peucker algorithm

PROCEDURE APPROXIMATE( $C, L, n, m, \theta$ ) outline


$C$  - contour points,  $L$ -line segments,  $n$  and  $m$  - point indices,  $\theta$  - distance threshold

- 1 Add a straight line  $l$  segment beginning in  $C(n)$  and ending in  $C(m)$  to the set of linear segments  $L$
- 2 For each point in  $C(i)$ , where  $n \leq i \leq m$  compute the distance  $d(C(i), l)$
- 3 Select  $k$  with the largest distance to the line segment  $d(C(k), l) = \max_{n \leq i \leq m} d(C(i), l)$
- 4 If for all computed  $d(C(k), l) \leq \theta$  then return from procedure
- 5 Remove  $l$  from the set  $L$
- 6 Call APPROXIMATE( $C, L, n, k, \theta$ ), APPROXIMATE( $C, L, k, m, \theta$ )

For a contour with point indices  $1 \dots N$  the procedure can be called as follows

APPROXIMATE( $C, L, 1, N, \theta$ )

# Approximating contour with a polyline




# Hough transform for line detection

Let us have straight line equation given in the form:

$$d = x \cos(\alpha) + y \sin(\alpha)$$

where

- $\alpha$  - the angle orientation of the vector normal to the line
- $d$  - the distance from the line to the coordinate system origin





# Hough transform for line detection

Then for every  $(x_i, y_i)$  pixel that is a part of the line the above equation must hold, so

$$d = x_i \cos(\alpha) + y_i \sin(\alpha)$$

Thinking the opposite way we will see that all lines going through the point  $(x_i, y_i)$  correspond to some sine-like function in a space  $[d, \alpha]$ . Let us now draw sine function for each point on the line in one graph.



Sines cross at the point  $[\alpha, d]!$

# Hough transform for line detection - basic algorithm

INPUT: image with detected contour elements

ALGORITHM outline:

- 1 Create a 2-dimensional array  $H$  (Hough accumulator) representing  $[d, \alpha]$  space. The space must be discretized within range

$$-\sqrt{M^2 + N^2} \leq d \leq \sqrt{M^2 + N^2}, -\frac{\pi}{2} \leq \alpha \leq \frac{\pi}{2}$$


where  $(M, N)$  is the image size

- 2 Set all elements of  $H$  to zero
- 3 For each contour element  $(x_i, y_i)$  and for all discrete values of  $\alpha^+$  do
  - compute  $d_i = x_i \cos(\alpha^+) + y_i \sin(\alpha^+)$
  - approximate  $(d_i, \alpha^+)$  with a nearest discrete value  $(d_i^+, \alpha^+)$
  - increment  $H[d_i^+, \alpha^+]$  by 1
- 4  $H[d, \alpha]$  contain votes for particular  $(d, \alpha)$  pairs. Maxima in the (smoothed) accumulator correspond to straight line segments.

# Hough transform for line detection - basic algorithm



(a)



(b)

(a) - Hough accumulator, (b) - input image



# Hough transform for line detection - faster algorithm

INPUT: image with detected contour elements and their orientations (normals)

ALGORITHM outline:

- 1 Create a 2-dimensional array  $H$  (Hough accumulator) representing  $[d, \alpha]$  space. The space must be discretized within range

$$-\sqrt{M^2 + N^2} \leq d \leq \sqrt{M^2 + N^2}, -\frac{\pi}{2} \leq \alpha \leq \frac{\pi}{2}$$

where  $(M, N)$  is the image size

- 2 Set all elements of  $H$  to zero
- 3 For each contour element  $(x_i, y_i)$  and for **established** contour element orientation  $\alpha_i$  do
  - compute  $d_i = x_i \cos(\alpha_i) + y_i \sin(\alpha_i)$
  - approximate  $(d_i, \alpha_i)$  with a nearest discrete value  $(d_i^+, \alpha_i^+)$
  - increment  $H[d_i^+, \alpha_i^+]$  by 1
  - $H[d, \alpha]$  contain votes for particular  $(d, \alpha)$  pairs. Maxima in the (smoothed) accumulator correspond to straight line segments.

**Exercise:** Why not use simpler parametrization (line-to-line conversion)

$$y = a \cdot x_i + b \Leftrightarrow b = -a \cdot x_i + y$$


# Hough transform for circle detection approach I

The circle  $[c_x, c_y, r]$  is given by the equation

$$r^2 = (x - x_c)^2 + (y - y_c)^2$$

Now, select a point on the circle boundary ( $e_i$ ) and construct a line going through this point and perpendicular to the circle boundary.

$g_i : y = a_i \cdot x + b_i$  for some boundary element  $e_i$



# Hough transform for circle detection approach I

The line obviously goes through the point  $(x_c, y_c)$ .

$$g_i : y_c = a_i \cdot x_c + b_i$$

Now, considering all such lines, the following condition must hold for their parameters  $[a, b]$

$$y_c = a \cdot x_c + b$$

which can be rewritten as


$$b = -x_c \cdot a + y_c$$

and  $b$  is a linear function of  $a$ . So for each boundary element  $(e_i)$  introduces a pair of line parameters  $(a_i, b_i)$ . The points  $(a_i, b_i)$  should lie on a single line (with unknown parameters  $(-x_c, y_c)$ ).

**Thus we just reduced a problem of detecting circles into a problem of detecting lines**

# Hough transform for circle detection approach I

Note: If we transform edge elements of real images of circles we will have multiple lines with imperfectly aligned points in  $[a, b]$  space.



# Hough transform for circle detection approach I

PROCEDURE outline:

- 1 Transform boundary elements ( $e_i$ ) from the image into  $[a, b]$  space
- 2 Perform line detection in  $[a, b]$  space (another Hough transform can be used here)
- 3 Restore  $(a_i, b_i)$  points assigned to a single detected straight line segment
- 4 Estimate  $(x_c, y_c)$  from the set of  $(a_i, b_i)$  points using LSE approach:

$$(c_x, c_y) = \underset{i}{\operatorname{argmin}} \sum \varepsilon_i^2(c_x, c_y)$$

$$\varepsilon_i(c_x, c_y) = b_i + x_c \cdot a_i - y_c$$

# Hough transform for circle detection approach II

A more direct Hough-transform-based approach to circles detection can be also given. The algorithm uses a 3-dimensional Hough accumulator  $H(x_c, y_c, r)$ . INPUT: image with detected boundary elements ( $e_i$ ) and their orientations (normals)  
PROCEDURE outline:

- ❶ Set all the elements of the Hough accumulator to 0  
( $H(x_c, y_c, r) = 0$ )
- ❷ For each boundary element  $e_i$  with coordinates  $(x_i, y_i)$  and contour normal direction  $\alpha$  do
  - For each allowed discrete value of circle radius  $r^+$  do
    - Compute circle center candidate  $(\tilde{x}_c, \tilde{y}_c)$ 
$$\tilde{x}_c = x_i + \cos(\alpha) * r^+$$
$$\tilde{y}_c = y_i + \sin(\alpha) * r^+$$
    - Round  $(\tilde{x}_c, \tilde{y}_c)$  to the nearest Hough accumulator grid coordinates  $(\tilde{x}_c^+, \tilde{y}_c^+)$  and increment by one  $H(\tilde{x}_c^+, \tilde{y}_c^+, r^+)$
- ❸  $H[x_c, y_c, r_c]$  contain votes for particular  $(x_c, y_c, r_c)$  triplets..  
Maxima in the (smoothed) accumulator correspond to detected circles.

# GHT for arbitrary contour detection


Generalized Hough Transform enables to detect arbitrary contour shapes in the image. The method is divided into two steps: Learning and Detection. For each detected contour, the following parameters are established:

- $\mathbf{x}_r = (x_r, y_r)$  - reference point of the contour
- $s$  - contour scale
- $\alpha$  - contour orientation

For now we will assume that scale  $s$  and  $\alpha$  are fixed.



# GHT for arbitrary contour detection



## Training - creating a lookup table R

- ❶ Select a reference point  $\mathbf{x}_R$  - e.g. a center of mass
- ❷ For each contour point  $\mathbf{x}_B$ 
  - compute  $\phi(\mathbf{x}_B)$  - gradient direction in  $\mathbf{x}_B$
  - compute  $\mathbf{r} = \mathbf{x}_R - \mathbf{x}_B$
  - store  $\mathbf{r}$  as a pair  $\phi_B \rightarrow \mathbf{r}_B$
- ❸ Discretize  $\phi$  and create R-table mapping gradient direction  $\phi$  into a sequence of possible  $\mathbf{r}$  vectors.

## Detection


- ❶ Determine all contour points  $\mathbf{x_I}$  together with their gradient orientations
- ❷ Initialize accumulator  $H = 0$
- ❸ For all contour points  $\mathbf{x_I}$  in the image do
  - Restore from R-table all values  $\mathbf{r}(\phi(\mathbf{x_I}))$  matching the current gradient  $\phi(\mathbf{x_I})$
  - For each  $\mathbf{r}(\phi(\mathbf{x_I}))$  increment by 1 a value of Hough accumulator cell  $\mathbf{x_I} + \mathbf{r}(\phi(\mathbf{x_I}))$  in a plane  $H(s, \alpha)$  of the Hough accumulator
- ❹ Smooth accumulator and detect local maxima
- ❺ Detected object reference points (with respective scales and orientations) are found in the positions of local maximas.

# GHT for arbitrary contour detection

GHT can be generalized for detecting object at different scales and orientations. This can be done by creating additional (or extending existing) R-tables

- 1 In order to prepare R-table for different scale, select scale  $s$  and multiply  $\mathbf{r}_B$  by  $s$  for each pair ( $\phi_b \rightarrow \mathbf{r}_b$ )
- 2 In order to prepare R-table for different angle, select rotation angle  $\alpha$  and add  $\alpha$  to  $\phi_B \pmod{2\pi}$  for each pair ( $\phi_b \rightarrow \mathbf{R}_\alpha \mathbf{r}_b$ ) before discretizing
- 3 In such way the original R-tables are converted to R-tables for different scales  $s$  and angles  $\alpha$
- 4 During detection for each contour elements all R-tables calculated for different  $s$  and  $\alpha$  must be applied. Therefore we must use an augmented Hough Accumulator  $H(x_b, y_b, s, \alpha)$

# GHT for arbitrary contour detection



# Random Sample Consensus (RanSaC)

RANSAC algorithm outline

INPUT:  $D$  - data points,  $\theta$  - maximum error of model fitting

Initialize inliers  $I = \emptyset$  and model  $M = null$

REPEAT until the maximum number of iterations is reached

- 1 Randomly select minimum number of data points from  $D$  needed to establish a model
- 2 Compute model parameters  $M_{current}$  basing on the selected points
- 3 Run through the data  $D$  and select all points matching the model  $M_{current}$  with error lower than  $\theta$  as  $I_{current}$
- 4 If  $|I_{current}| > |I|$ , then  $I = I_{current}$ ,  $M = M_{current}$

return  $M$  and  $I$

# Random Sample Consensus (RanSaC)

## Remarks:

- The algorithm may take into account not only the number inliers but also overall quality of model fit
- The model can be filtered after generation (e.g. to accept only models with certain parameter values)
- The required number of iterations can be established basing on a pre-defined *confidence level*  $p$ . Confidence level is  $p$  if the probability of finding true object by RanSaC is greater than  $p$ .

# Random Sample Consensus (RanSaC)

Using confidence  $p$  to obtain number of iterations

- Let  $N$  be the number of points in the set and  $N_o$  the number of points in our object. Let  $w = \frac{N_o}{N}$  be the probability of selecting inlier in a one draw.
- The probability of selecting enough inliers to estimate the model is  $w^n$ , where  $n$  is the number of points required to establish the model
- The probability of successful establishment of at least one valid model in  $T$  iterations is

$$1 - (1 - w^n)^T = p$$

- And the minimum number of iterations to obtain confidence  $p$

$$T \geq \frac{\log(1 - p)}{\log(1 - w^n)}$$

- Number of points in the object  $N_o$  is typically unknown, but can be approximated by  $|I|$  during the algorithm execution



# Random Sample Consensus (RanSaC)

After finishing the RanSaC procedure the resulting set of inliers can be iteratively adjusted according to the procedure REPEAT until the set  $I$  stabilizes

- 1 Reestimate model parameters  $\tilde{M}$  using ALL inliers from  $I$
- 2 Run through the data  $D$  and select all points matching the model  $\tilde{M}$  with an error lower than  $\theta$  as  $I$

# Random Sample Consensus (RanSaC)

RanSaC is a generic parametric model-fitting tool that is robust to outliers. Some applications in Computer Vision include

- Detection of objects in 2D and 3D images
- Estimation of image transformations (e.g. homography, fundamental or essential matrices) basing on noisy point correspondences
- Fitting curves and surfaces to noisy data

# Feature description

Abundance of feature descriptors depending on the type of feature used, e.g.:


- Geometric properties (geometric moments) or Zernike's moments for geometric description
- Various pixel or color histograms
- Contour descriptors for contour features
- **SIFT**, SURF, ORB, HoG, wavelet (Haar, Gabor) descriptors for image patches
- Bag of Words features
- ...

# Scale Invariant Features Descriptor

Assume that we already have SIFT (or other type of) keypoints. For each keypoint using the image scale closest to the scale of detected keypoint

- 1 a 16x16 grid is imposed on the image in keypoint position (coordinates of the descriptor and gradient location are rotated according to keypoint orientation)
- 2 for each 4x4 cell a 8-bin histogram of gradient orientations is computed (Gaussian weighing with  $\sigma$  corresponding to descriptor window size is used to emphasize center pixels)
- 3 the descriptor is formed from all histogram bins (128 elements)
- 4 descriptor is normalized to unit length and its particular values are thresholded

# Scale Invariant Features Descriptor



SIFT descriptor 8x8 version

Source: David. G. Lowe, Distinctive image features from scale-invariant keypoints, 2004

# Automatic feature detection and description

# Automatic feature detection and description


- Machine Learning algorithms may be used to craft feature detection algorithms
- Some of them base on unsupervised learning - generic feature detectors (e.g. clustering methods)
- Others base on supervised learning - specialized feature detectors

These are the **Feature Learning** methods.



# Convolutional Neural Networks for feature detection

- CNNs inherently use feature detection in their convolutional layers
- Each convolutional layer takes as an input image or feature map and outputs another set of feature maps (interpreted as feature densities), higher layers incorporate information from larger areas
- Higher layers aggregate larger portions of image
- Computed feature maps are not always human-interpretable (features are generated to serve some purpose - e.g. recognition of some object class), but frequently can be re-used.




Source: [https://www.quora.com/](https://www.quora.com/How-is-a-convolutional-neural-network-able-to-learn-invariant-features)

How-is-a-convolutional-neural-network-able-to-learn-invariant-features

# Convolutional Neural Networks for feature detection

## Sample features



Source: <https://www.deeplearningbook.org>

# Specialized feature detectors


Does human brain use specialized feature detectors?



Source: <https://www.businessinsider.com/>

# Specialized feature detectors


Does human brain use specialized feature detectors?



Source: <https://www.businessinsider.com/>

# Specialized feature detectors

## Deep Dream



Source: <https://en.wikipedia.org/wiki/DeepDream>

- Dedicated solutions to keypoint detection/feature extraction are also present

# Exercises

# Exercise 1. Hough transform for line detection

The following table represents pixel coordinates and orientations for contours detected in the image:

$x_i$	$y_i$	$r(\alpha_i)$
1	1	3
2	2	3
3	3	3
4	3	3
5	5	3
3	2	6
4	2	6
5	2	6
7	3	6
1	4	0
2	4	0
3	4	0

- Assume that gradient directions are discretized into 8 bins, with  $r(0) = 0$  and  $r(45) = 1 \dots$
- Apply the Hough transform for line segment detection
- Assume that at least 2 contour elements must vote for each line segment
- Provide the resulting Hough accumulator, intermediate results and an illustration



## Exercise 2. Hough transform for circle detection

The following table represents pixel coordinates and orientations for contours detected in the image:

$x_i$	$y_i$	$e(\alpha_i)$
1	0	1
5	0	3
1	4	7
5	4	5

- Assume that gradient directions are discretized into 8 bins, with  $e(0^\circ) = 0$  and  $e(45^\circ) = 1 \dots$
- Apply the Hough transform 'approach I' for circle detection
- Provide provide intermediate results and relevant illustrations

## Exercise 3. Generalized Hough transform

You are given a template of a white  $6 \times 4$  rectangle with rounded corners on a black background. Some points on the perimeter of the rectangle were sampled:


- $(3,0)$ ,  $(3,2)$ ,  $(0,2)$ ,  $(-3,2)$ ,  $(-3,0)$ ,  $(-3,-2)$ ,  $(0,-2)$ ,  $(3,-2)$

Assume that the reference point is in the mass centre of the rectangle. Gradient direction in the corner is the average of the gradient directions of the adjacent sides.

- 1 Synthesize GHT R-table for this sample shape
- 2 How many points are there in each R-table entry? What if we add additional contour element  $(1,2)$ ?
- 3 Synthesize GHT R-table for this shape rotated by 45 degrees

## Exercise 3. Generalized Hough transform

There is given an image with detected contour points and orientations (orientations are discretized into 8 bins, shading does not reflect the actual colors).



- 1 Apply the GHT to find a single reference point in the detection of the sample rectangle. Use the rotated version of R-table.
- 2 Did we obtain one or multiple solutions?
- 3 Provide illustrations and intermediate results.

## Exercise 4. RanSaC

You perform ellipse detection in the image using RanSaC algorithm. The ellipse equation is given as

$$\frac{(x - x_0)^2}{a^2} + \frac{(y - y_0)^2}{b^2} = 1$$

You detected 50000 contour points and you expect the detected ellipse to contain at least 10000 points. How many RanSaC iterations must be performed in order to detect an ellipse with the confidence of 0.99?

## Exercise 5. Douglas-Peucker algorithm

Apply the Douglas-Peucker algorithm with distance threshold  $\theta = 2$  to the polyline with nodes

- $(2, 2) - (4, 5) - (6, 1) - (12, 3) - (14, 2)$

Provide intermediate states of the polyline.

# Exercise 1. Hough transform for line detection - hints

The following table represents pixel coordinates and orientations for contours detected in the image:

$x_i$	$y_i$	$r(\alpha_i)$
1	1	3
2	2	3
3	3	3
4	3	3
5	5	3
3	2	6
4	2	6
5	2	6
7	3	6
1	4	0
2	4	0
3	4	0

- Assume that gradient directions are discretized into 8 bins, with  $r(0) = 0$  and  $r(45) = 1 \dots$
- Apply the Hough transform for line segment detection
- Assume that at least 2 contour elements must vote for each line segment
- Provide the resulting Hough accumulator, intermediate results and an illustration

# Exercise 1. Hough transform for line detection - hints

Extended table:  $d_i = x_i \cos(\alpha) + y_i \sin(\alpha)$

$x_i$	$y_i$	$r(\alpha_i)$	$\sin(\alpha_i)$	$\cos(\alpha_i)$	$d_i$	$d_i^+$
1	1	3	$\frac{\sqrt{2}}{2}$	$-\frac{\sqrt{2}}{2}$	0	0
2	2	3	$\frac{\sqrt{2}}{2}$	$-\frac{\sqrt{2}}{2}$	0	0
3	3	3	$\frac{\sqrt{2}}{2}$	$-\frac{\sqrt{2}}{2}$	0	0
4	3	3	$\frac{\sqrt{2}}{2}$	$-\frac{\sqrt{2}}{2}$	$-4\frac{\sqrt{2}}{2} + 3\frac{\sqrt{2}}{2} = -\frac{\sqrt{2}}{2}$	-1
5	5	3	$\frac{\sqrt{2}}{2}$	$-\frac{\sqrt{2}}{2}$		0
3	2	6	-1	0	-2	-2
4	2	6	-1	0	-2	-2
5	2	6	-1	0	-2	-2
7	3	6	-1	0	-3	-3
1	4	0	0	1	1	1
2	4	0	0	1	2	2
3	4	0	0	1	3	3

# Exercise 1. Hough transform for line detection - hints

Hough accumulator (red - local maxima of votes)

$d_i^+ / r(\alpha_i)$	0	1	2	3	4	5	6	7
-3							1	
-2							3	
-1				1				
0				4				
1	1							
2	1							
3	1							

Line 1:  $x \cos(135^\circ) + y \sin(135^\circ) = 0 \rightarrow y = x$

Line 2:  $x \cos(270^\circ) + y \sin(270^\circ) = -2 \rightarrow y = 2$



## Exercise 2. Hough transform for circle detection - hints

The following table represents pixel coordinates and orientations for contours detected in the image:

$x_i$	$y_i$	$e(\alpha_i)$
1	0	1
5	0	3
1	4	7
5	4	5

- Assume that gradient directions are discretized into 8 bins, with  $e(0^\circ) = 0$  and  $e(45^\circ) = 1 \dots$
- Apply the Hough transform 'approach I' for circle detection
- Provide provide intermediate results and relevant illustrations

## Exercise 2. Hough transform for circle detection - hints

Basing on contour points and their orientation one can find lines normal to contour elements. Each line is described by the equation  $y = ax + b$

$x_i$	$y_i$	$e(\alpha_i)$	$\alpha_i$	$a_i = \tan(\alpha_i)$	$b_i$
1	0	1	$45^\circ$	1	-1
5	0	3	$135^\circ$	-1	5
1	4	7	$315^\circ$	-1	5
5	4	5	$225^\circ$	1	-1

It can be observed that pairs of coordinates in the Hough space  $(a_i, b_i)$  are compressed to just 2 points  $(1, -1), (-1, 5)$ . Through these points we can (in the Hough space) draw a line  $b = -x_0a + y_0$ . In our case the line has equation  $b = -3a + 2$ . Therefore the detected circle center has coordinates  $(x_0, y_0) = (3, 2)$ .

**Remark:** For larger number of points  $(a_i, b_i)$  in the Hough space it is necessary to solve the regression problem.

## Exercise 3. Generalized Hough transform - hints

You are given a template of a white  $6 \times 4$  rectangle with rounded corners on a black background. Some points on the perimeter of the rectangle were sampled:


- $(3,0)$ ,  $(3,2)$ ,  $(0,2)$ ,  $(-3,2)$ ,  $(-3,0)$ ,  $(-3,-2)$ ,  $(0,-2)$ ,  $(3,-2)$

Assume that the reference point is in the mass centre of the rectangle. Gradient direction in the corner is the average of the gradient directions of the adjacent sides.

- 1 Synthesize GHT R-table for this sample shape
- 2 How many points are there in each R-table entry? What if we add additional contour element  $(1,2)$ ?
- 3 Synthesize GHT R-table for this shape rotated by 45 degrees

## Exercise 3. Generalized Hough transform - hints

There is given an image with detected contour points and orientations (orientations are discretized into 8 bins, shading does not reflect the actual colors).



- 1 Apply the GHT to find a single reference point in the detection of the sample rectangle. Use the rotated version of R-table.
- 2 Did we obtain one or multiple solutions?
- 3 Provide illustrations and intermediate results.

# Exercise 3. Generalized Hough Transform - hints

ad. 1

- Rounded corners - we assume that in rectangle corners we have edge directions : 1,3,5,7
- For example, in  $(2, -3)$ , we have direction 3 and a vector  $\mathbf{r}(3) = [-3, 2]^T$

$\phi_i$	$\mathbf{r}(\phi_i)$
0	$[3, 0]^T$
1	$[3, 2]^T$
2	$[0, 2]^T$
3	$[-3, 2]^T$
4	$[-3, 0]^T$
5	$[-3, -2]^T$
6	$[0, -2]^T$
7	$[3, -2]^T$

## Exercise 3. Generalized Hough Transform - hints

ad. 2

- After adding a new contour element (1, 2), R-table for direction 6 is extended by a new vector.

$\phi_i$	$\mathbf{r}(\phi_i)$
0	$[3, 0]^T$
1	$[3, 2]^T$
2	$[0, 2]^T$
3	$[-3, 2]^T$
4	$[-3, 0]^T$
5	$[-3, -2]^T$
6	$[0, -2]^T, [-1, -2]^T$
7	$[3, -2]^T$

# Exercise 3. Generalized Hough Transform - hints

ad. 2

- Rotation  $\alpha = \pi/4$
- All elements in R-table are first circularly shifted by 1

- Rotation matrix:

$$\mathbf{R}_\alpha = \begin{bmatrix} \frac{\sqrt{2}}{2} & -\frac{\sqrt{2}}{2} \\ \frac{\sqrt{2}}{2} & \frac{\sqrt{2}}{2} \end{bmatrix}$$

- Rotation formula:


$$\begin{aligned} \mathbf{R}_\alpha \mathbf{x} &= \frac{\sqrt{2}}{2} \begin{bmatrix} 1 & -1 \\ 1 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} \\ &= \frac{\sqrt{2}}{2} [x - y, x + y]^T \end{aligned}$$

$\phi_i$	$\mathbf{r}(\phi_i)$
0	$\frac{\sqrt{2}}{2} [3 - (-2), 3 + (-2)]^T = \frac{\sqrt{2}}{2} [5, 1]^T \approx [4, 1]^T$
1	$\frac{\sqrt{2}}{2} [3 - 0, 3 + 0]^T = \frac{\sqrt{2}}{2} [3, 3]^T \approx [2, 2]^T$
2	$\frac{\sqrt{2}}{2} [3 - 2, 3 + 2]^T = \frac{\sqrt{2}}{2} [1, 5]^T \approx [1, 4]^T$
3	$[-1, 1]^T$
4	$[-4, -1]^T$
5	$[-2, -2]^T$
6	$[-1, -4]^T$
7	$[1, -1]^T$

# Exercise 3. Generalized Hough Transform - hints

ad. 2

- dots - vote counts
- blue - sample votes in the Hough accumulator
- red - the only maximum (single detection of the rectangle)





## Exercise 4. RanSaC - hints

You perform ellipse detection in the image using RanSaC algorithm. The ellipse equation is given as

$$\frac{(x - x_0)^2}{a^2} + \frac{(y - y_0)^2}{b^2} = 1$$

You detected 50000 contour points and you expect the detected ellipse to contain at least 10000 points. How many RanSaC iterations must be performed in order to detect an ellipse with the confidence of 0.99?

- Number of points required to estimate the model of ellipse is 4
- Required number of iterations is 2876

## Exercise 5. Douglas-Peucker algorithm - hints

### Douglas-Peucker

Apply the Douglas-Peucker algorithm with distance threshold  $\theta = 2$  to the polyline with nodes

- $(2, 2) - (4, 5) - (6, 1) - (12, 3) - (14, 2)$

Provide intermediate states of the polyline.

## Exercise 5. Douglas-Peucker algorithm - hints

The state of the estimated polyline in subsequent iterations of the algorithm.

- ❶  $(2, 2)-(14, 2)$
- ❷  $(2, 2)-(4, 5)-(14, 2)$
- ❸  $(2, 2)-(4, 5)-(6, 1)-(14, 2)$